


# Artificial Neural Networks


Wulfram Gerstner

EPFL, Lausanne, Switzerland


# The brain: Cortical Areas



frontal cortex




# The brain: Cortical Areas




# The Brain: zooming in


1mm



10 000 neurons  
3 km of wire




*Ramon y Cajal*




# The brain: a network of neurons

1mm




10 000 neurones  
3 km de cablage



*Ramon y Cajal*


Signal:

Potentiel d'action (impulsion)




# The brain: signal transmission

Signal:  
action potential (short pulse)




plus que mille entrées

# The brain: neurons sum their inputs





# Summary: the brain is a large network of neurons






● Active neuron




# Learning in the brain: changes between connections



**Neurons**



**Synapse**



**learning = change of connection**


# Artificial Neural Networks

Wulfram Gerstner

EPFL, Lausanne, Switzerland


1. The brain
2. Artificial Neural Networks

# Modeling: artificial neurons




- responses are added
- pulses created at threshold
- transmitted to other


→ Mathematical description



# Modeling: artificial neurons





forget spikes: continuous activity  $x$   
forget time: discrete updates



# Learning of connections in biology

*Where do the connection weights come from?*




## Hebbian Learning

When an axon of cell **j** repeatedly or persistently takes part in firing cell **i**, then **j**'s efficiency as one of the cells firing **i** is increased

*Hebb, 1949*


- local rule
- simultaneously active neurons

# Hebbian Learning of Associations






# Hebbian Learning of Associations




item memorized by change of synaptic weights

# Hebbian Learning: Associative Recall

**Recall:  
Partial info**




item recalled






# Neurons and Synapses form a big network



Brain



1mm

10 000 neurons

3 km of wire

10 billions neurons

10 000 connexions/neurons

**memory in the connections**


Distributed Architecture

**No separation of  
processing and memory**

# Quiz: biological neural networks


- ☐ Neurons in the brain have a threshold.
- ☐ Learning means a change in the threshold.
- ☐ Learning means a change of the connection weights
- ☐ The total input to a neuron is the weighted sum of individual inputs
- ☐ The neuronal network in the brain has no recurrent connections and no feedback connections

# Artificial Neural Networks for classification



# Artificial Neural Networks for classification

**Aim of learning:**  
Adjust connections such  
that output is correct  
(for each input)



# Artificial Neural Networks

Wulfram Gerstner

EPFL, Lausanne, Switzerland

1. The brain
2. Artificial Neural Networks
  - for classification
  - for action learning




# Artificial Neural Networks for action learning



Coactivation of 2 neurones:

- Connections strengthened
- Facilitates to repeat same action

**Even the mistakes?**




Missing:

Value of action

- 'goodie' for do
- 'success'
- 'compliment'



# Modeling – the role of reward



## Three factors for changing a connection


- activity of neuron  $j$
- activity of neurone  $i$
- success

*Barto 1985, Schultz et al. 1997; Waelti et al., 2001;  
Reynolds and Wickens 2002;  
Lisman et al. 2011*



# Deep reinforcement learning


## Chess



Artificial neural network  
(*AlphaZero*) discovers different  
strategies by playing against itself.

In Go, it beats Lee Sedol

## Go





# Deep reinforcement learning

## Network for choosing action


action:

*Advance king*

2<sup>e</sup> output for value of action:

*probability to win*

output



**learning:**

- change connections

**aim:**

- Predict value of position
- Choose next action to win




# Deep reinforcement learning (alpha zero)

*Silver et al. (2017) , Deep Mind*

output: 4672 actions

**Training** 44Mio  
games (9 hours)

*advance king*



input: 64x6x8x2 neuronss  
(about 10 000)

**Planning:**  
**potential sequences**  
(during 1s before playing  
next action)



# Deep reinforcement learning (alpha zero)



## Chess:

-discovers classic openings



-beats best human players

-beats best classic AI algorithms



A10: English Opening





D06: Queens Gambit





A46: Queens Pawn Game





E00: Queens Pawn Game




E61: Kings Indian Defence




C00: French Defence



B50: Sicilian Defence



B30: Sicilian Defence






# Self-driving cars


<https://selfdrivingcars.mit.edu/>

*Lex Friedman, MIT*

*advance and accerate*



Value: security,  
duration of travel



**External**

- 1. Radar
- 2. Visible-light camera
- 3. LIDAR
- 4. Infrared camera
- 5. Stereo vision
- 6. GPS/IMU
- 7. CAN
- 8. Audio

**Internal**

- 1. Visible-light camera
- 2. Infrared camera
- 3. Audio



Road Overlay:

Safety System ⬆⬇⬇⬆

# Artificial Neural Networks

Wulfram Gerstner


EPFL, Lausanne, Switzerland

1. The brain
2. Artificial Neural Networks
  - for classification
  - for action learning
  - for sequences (music, translation, speech)



# Deep networks with recurrent connections

*‘a man sitting on a couch with a dog’*



Network describes the  
image with the words:

*‘a man sitting on a couch with a dog’*

(Fang et al. 2015)

# Quiz: Classification versus Reinforcement Learning

- ☐ Classification is based on rewards
- ☐ Reinforcement learning is based on rewards
- ☐ Reinforcement learning aims at optimal action choices
- ☐ Classification aims at predicting the correct category such as 'car' or 'dog'

# Artificial Neural Networks

Wulfram Gerstner

EPFL, Lausanne, Switzerland

1. The brain
2. Artificial Neural Networks
  - for classification
  - for action learning
  - for sequences
3. Overview of class



# Artificial Neural Networks

Wulfram Gerstner




EPFL, Lausanne, Switzerland

1. Simple perceptrons for classification
2. Backprop and multilayer perceptron
3. Statistical Classification by deep networks
4. Deep learning:  
    regularization and tricks of the trade
5. Complements to deep learning
6. Sequence predictions and LSTMs
7. Convolutional networks
8. Reinforcement learning1: TD learning
9. Reinforcement learning2: SARSA
10. Reinforcement learning3: Policy Gradient
11. Deep Reinforcement learning
12. Applications

# Artificial Neural Networks

Wulfram Gerstner

EPFL, Lausanne, Switzerland

1. Simple perceptrons for classification
2. Backprop and multilayer perceptron
3. Statistical Classification by deep networks  [miniproject1](#)
4. Deep learning:  
    regularization and tricks of the trade
5. Complements to deep learning
6. Sequence predictions and LSTMs  [miniproject2](#)
7. Convolutional networks
8. Reinforcement learning1: TD learning
9. Reinforcement learning2: SARSA
10. Reinforcement learning3: Policy Gradient  [miniproject3](#)
11. Deep Reinforcement learning
12. Applications

Miniprojects (MPs): we use software package 'Keras'

- hand in 2 (or 3) out of 3 projects
- graded on a scale of 1-6
- average grade of MPs counts  $\frac{1}{3}$  toward final grade
- we do fraud detection interviews
- you get three weeks for each MP
- MP can be done alone or in group of two students
- interview for final MP is in first week after end of classes

→ **plan ahead!!**

**Written exam:**

- counts  $\frac{2}{3}$  toward final grade
- no tools allowed (no calculator, no cell phone, no paper, no book)
- 'mathy'

# Exercise sessions as follows:

- hand-out of exercise sheet  $n$  Friday of week  $n$
- You work on it until Thursday of week  $n+1$
- Solutions posted at noon, Thursday  $n+1$
- Friday week  $n+1$  by 10:00 am: you vote for the one exercise that you want to be explained
- Friday week  $n+1$  at 12:00 am. This exercise is explained on the blackboard + Q&A  
+ Q&A for miniprojects

TA's: - Dr. Johannni Brea  
- Dane Corneil  
- Florian Colombo  
- Teo Stocco

# Artificial Neural Networks

Wulfram Gerstner

EPFL, Lausanne, Switzerland

- The math is developed on the blackboard
- There are no written course notes!!
- All of the contents are standard textbook material

Choose a textbook that you like:

For first half of class:

- *Pattern Recognition and Machine Learning*, C.M Bishop, 2006
- *Neural Networks for Pattern Recognition*, C.M. Bishop, 1995
- *Deep Learning*, Ian Goodfellow et al., 2017 (also online)

For second half of class:

- *Reinforcement learning*, R. Sutton+ A. Barto (2<sup>nd</sup> ed, online)

# Artificial Neural Networks

## **Prerequisites:**

CS433, Pattern Classification and Machine Learning  
(Profs Jaggi+Urbanke)

## **Rules:**

If you have taken this class: please ask many questions

If you have not taken this class: please do not complain

# Artificial Neural Networks

## **Learning outcomes:**

- apply learning in deep networks to real data
- assess/evaluate performance of learning algorithms
- Elaborate relations between different mathematical concepts of learning
- judge limitations of learning algorithms
- propose models for learning in deep networks

## **Transversal skills:**

Access and evaluate appropriate sources of information

Manage priorities

work through difficulties, write a technical report

# Artificial Neural Networks

## **Work load:**

**4 credit course → 6 hours per week for 18 weeks**

**(1 ECTS = 27 hours of work)**



# Questions?

*... before we start*

TA's this year:

- Dr. Johannni Brea
- Dane Corneil
- Florian Colombo
- Teo Stocco

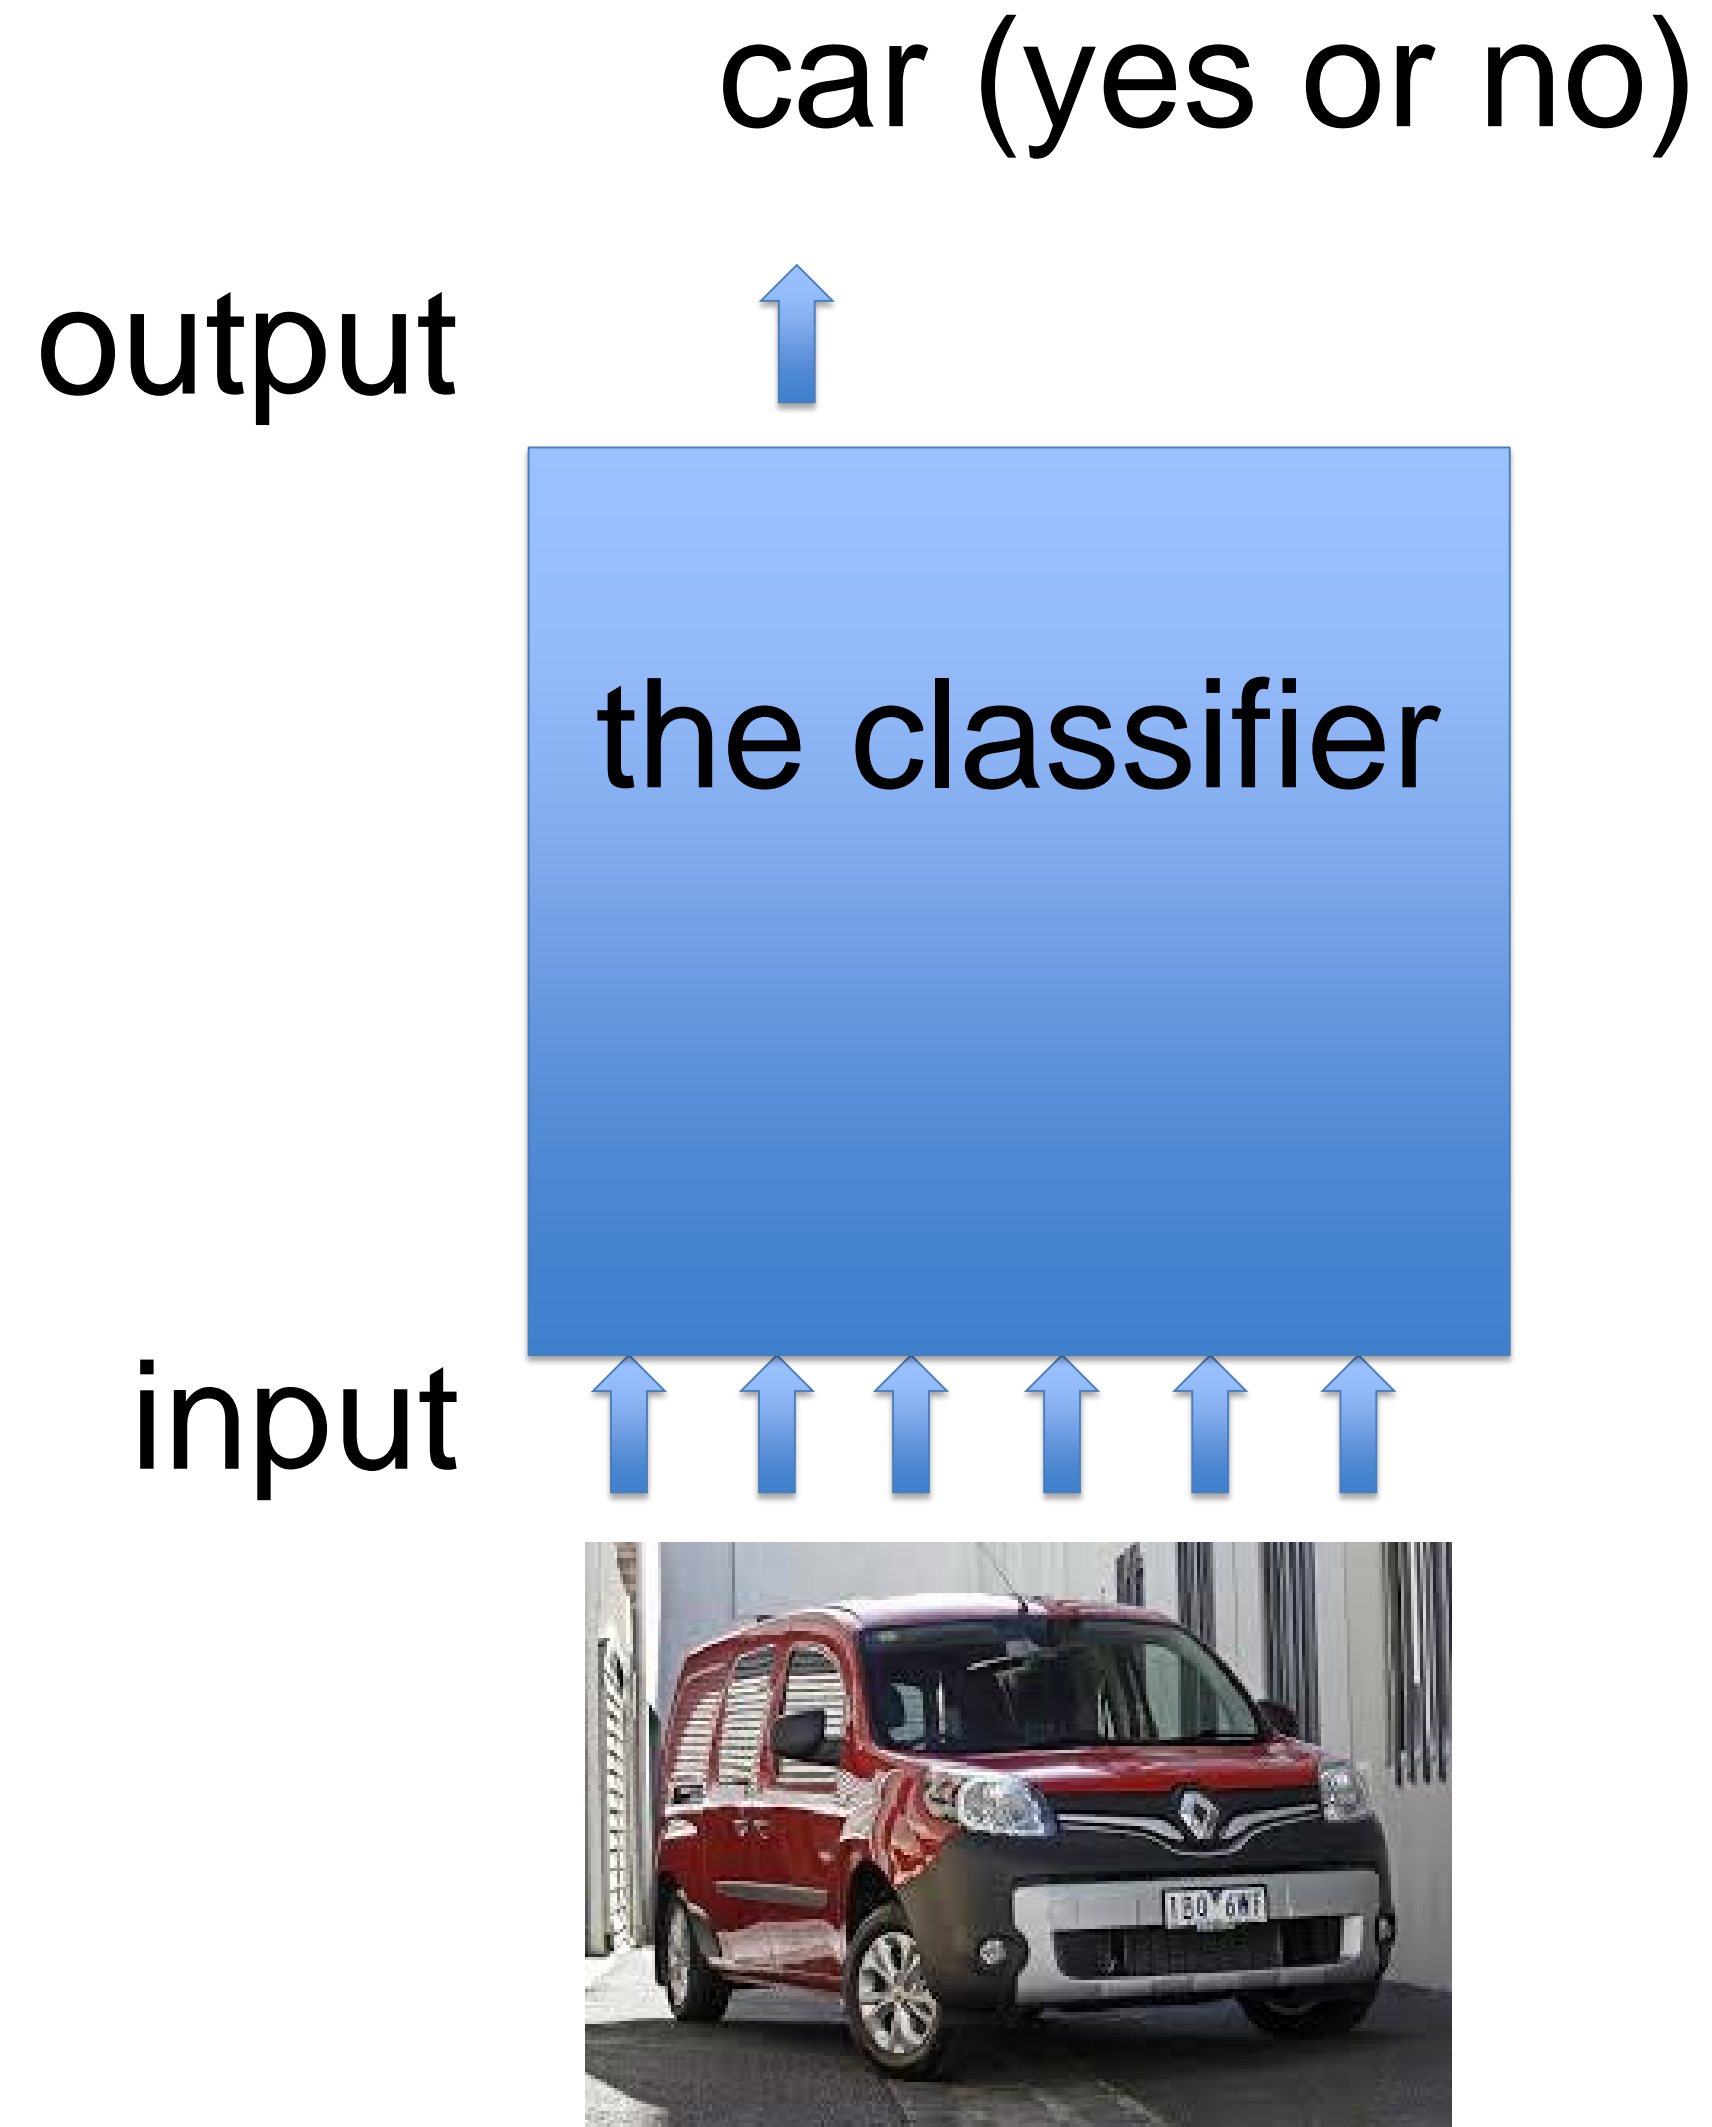
# Artificial Neural Networks: Lecture 1

## Simple Perceptrons for Classification

### **Objectives for today:**


- understand classification as a geometrical problem
- discriminant function of classification
- linear versus nonlinear discriminant function
- perceptron algorithm
- gradient descent for simple perceptrons

# 1. The problem of Classification



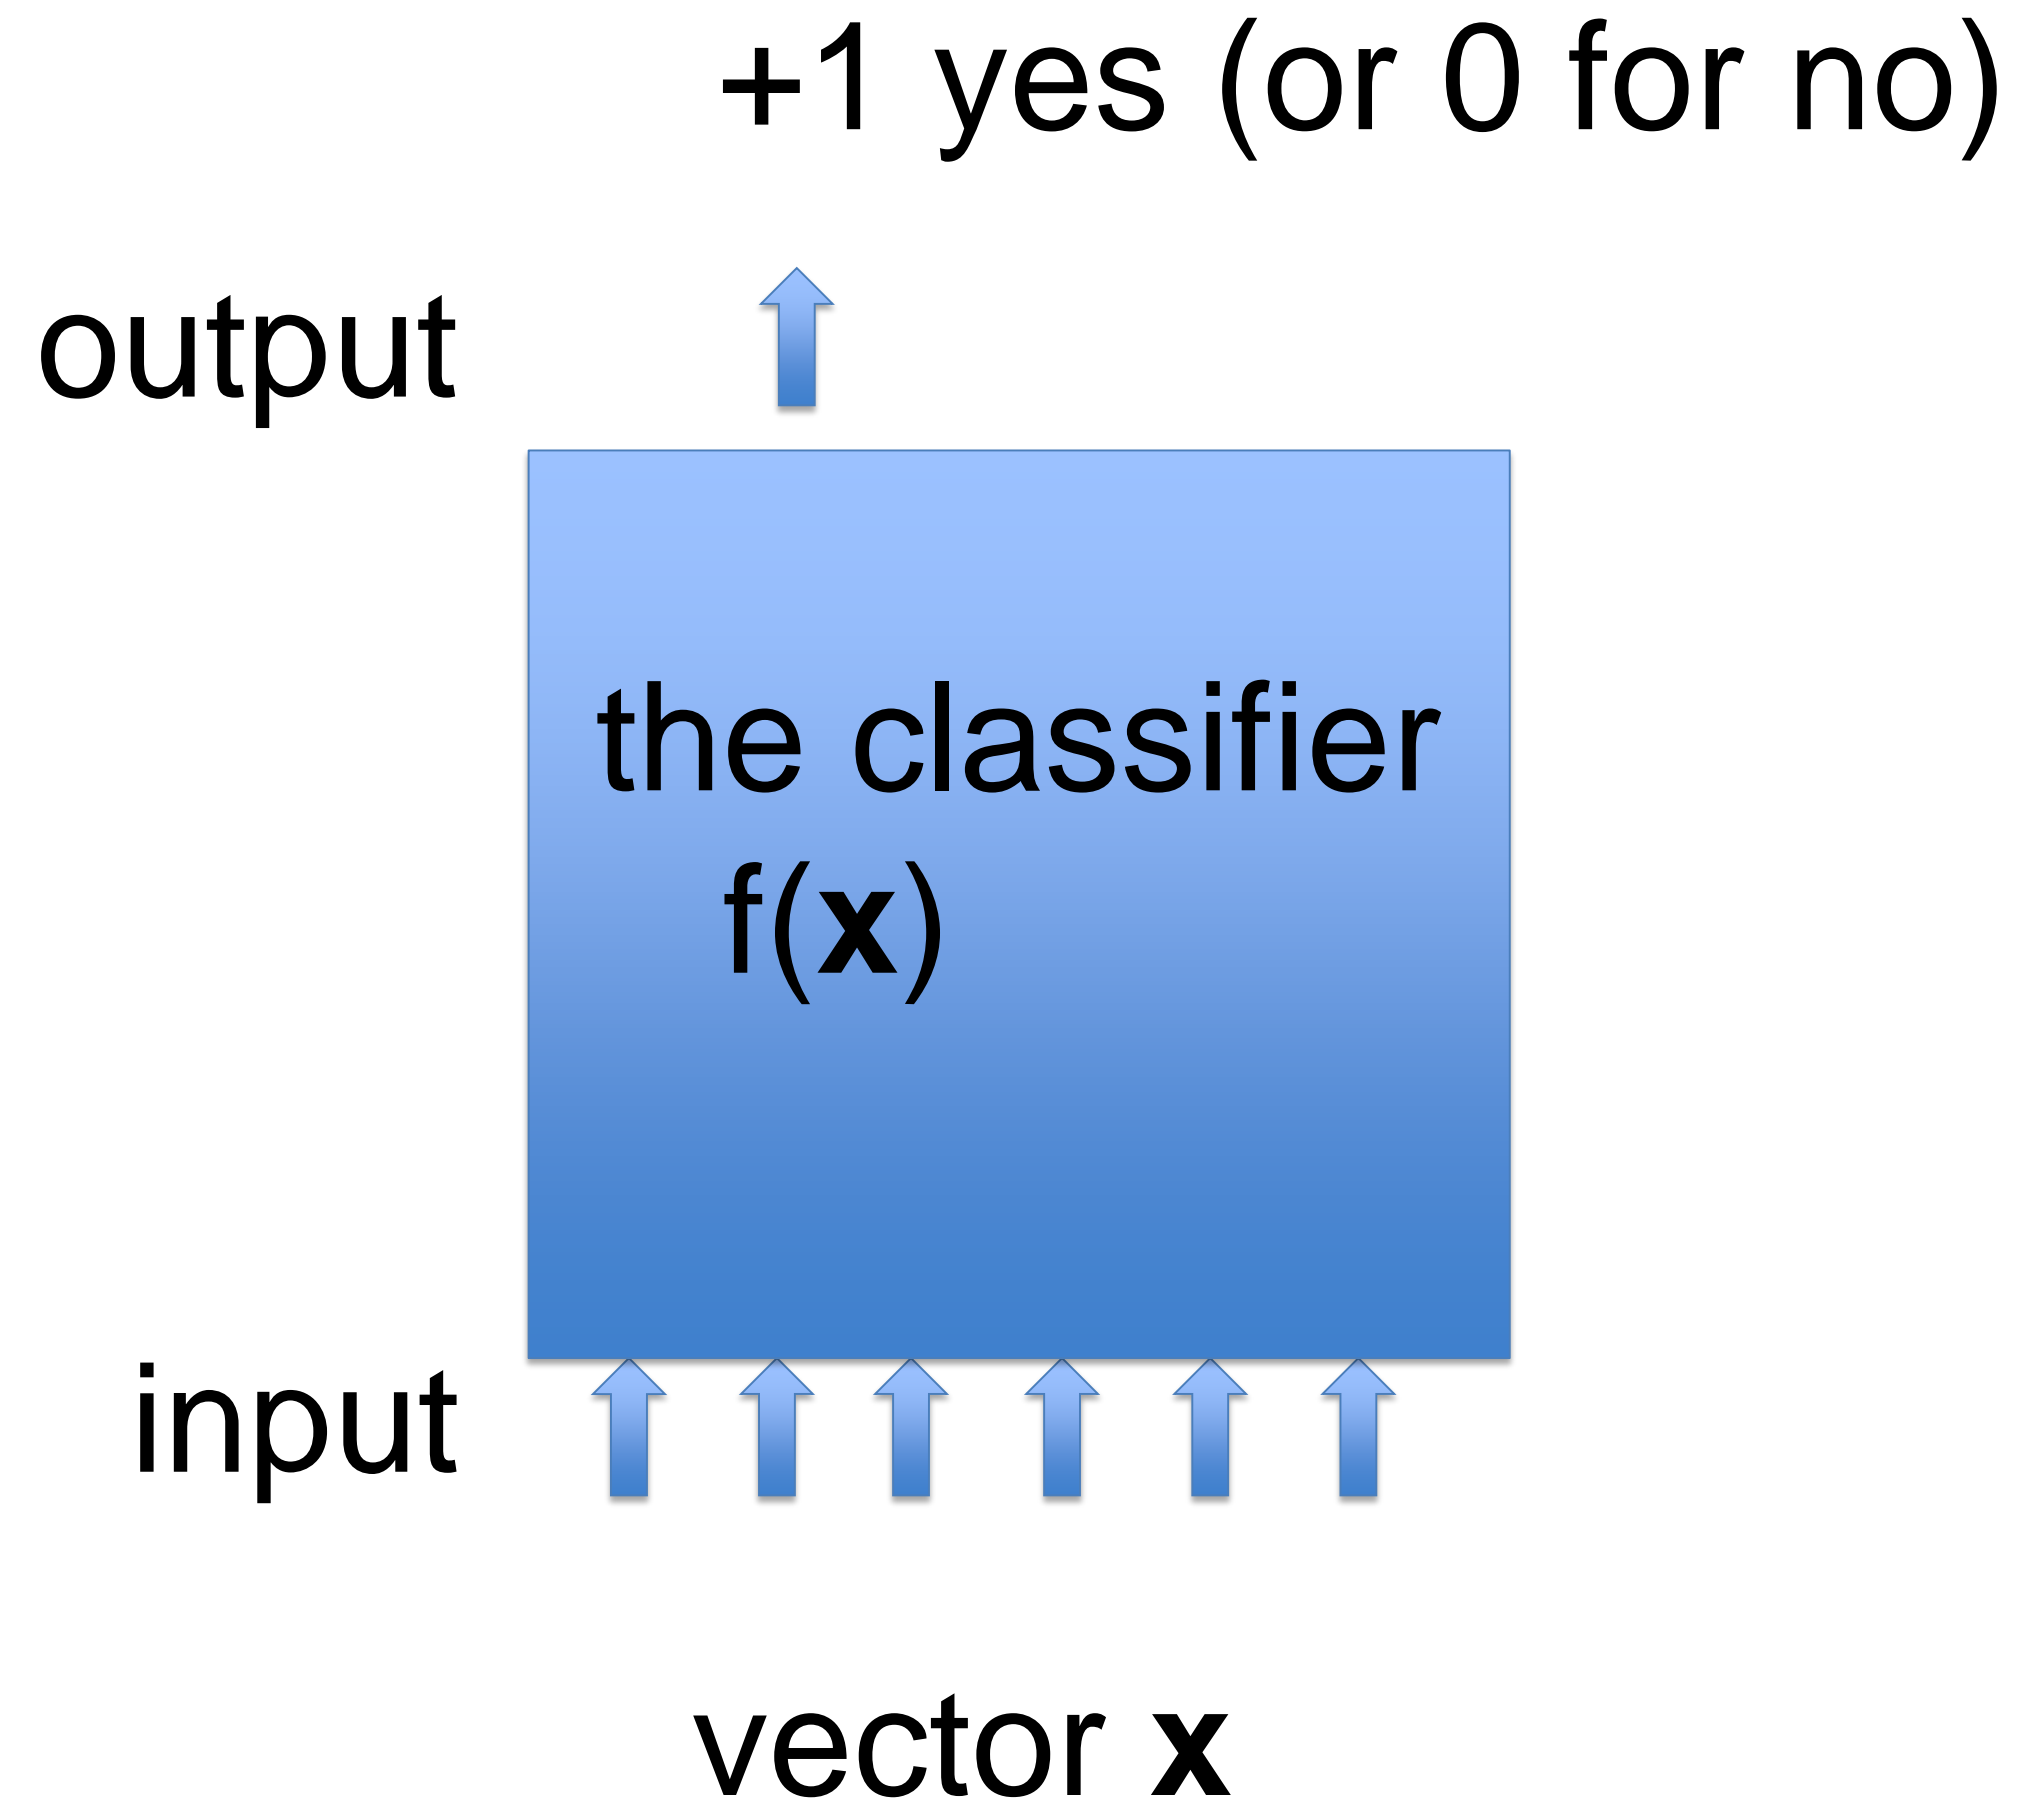
# 1. The problem of Classification

*Blackboard 1:*  
from images to vector




*Blackboard 1:*  
from images to vector

# 1. The problem of Classification





# 1. Classification as a geometric problem



*Blackboard 2:*  
from vectors to classification

*Blackboard 2:*  
from vectors to classification





# 1. Classification as a geometric problem

## Task of Classification

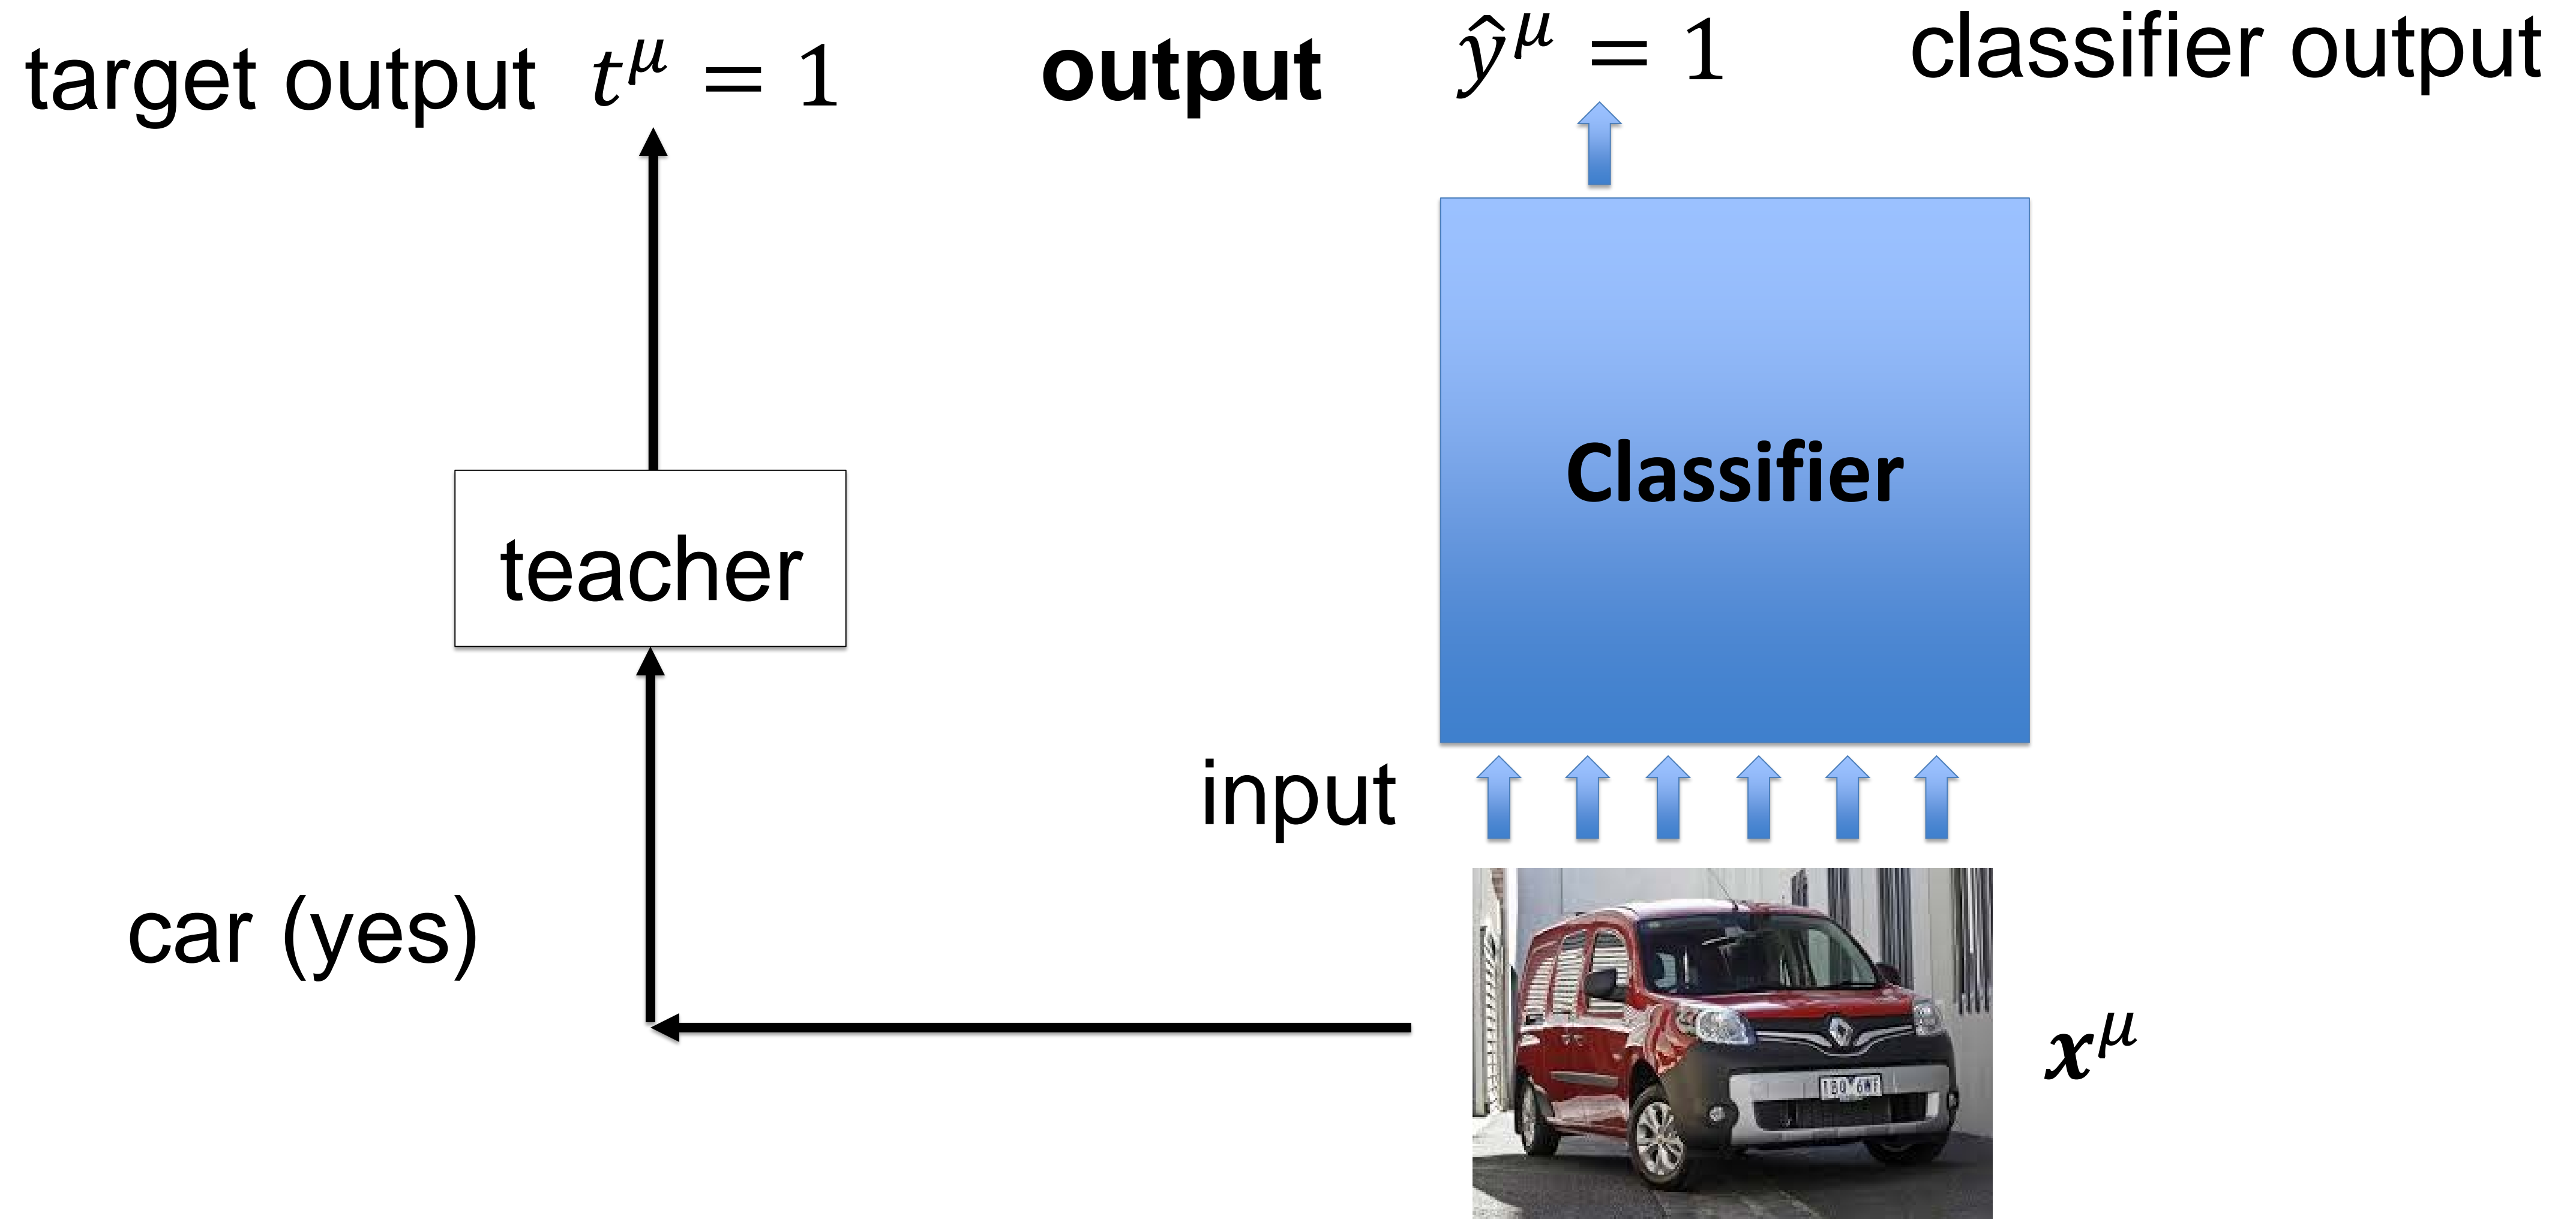
= find a **separating surface** in the high-dimensional input space

Classification by **discriminant function**  $d(\mathbf{x})$

→  $d(\mathbf{x})=0$  on this surface;  $d(\mathbf{x})>0$  for all positive examples  $\mathbf{x}$   
 $d(\mathbf{x})<0$  for all counter examples  $\mathbf{x}$



## 2. Data base for Supervised learning



## 2. Data base for Supervised learning

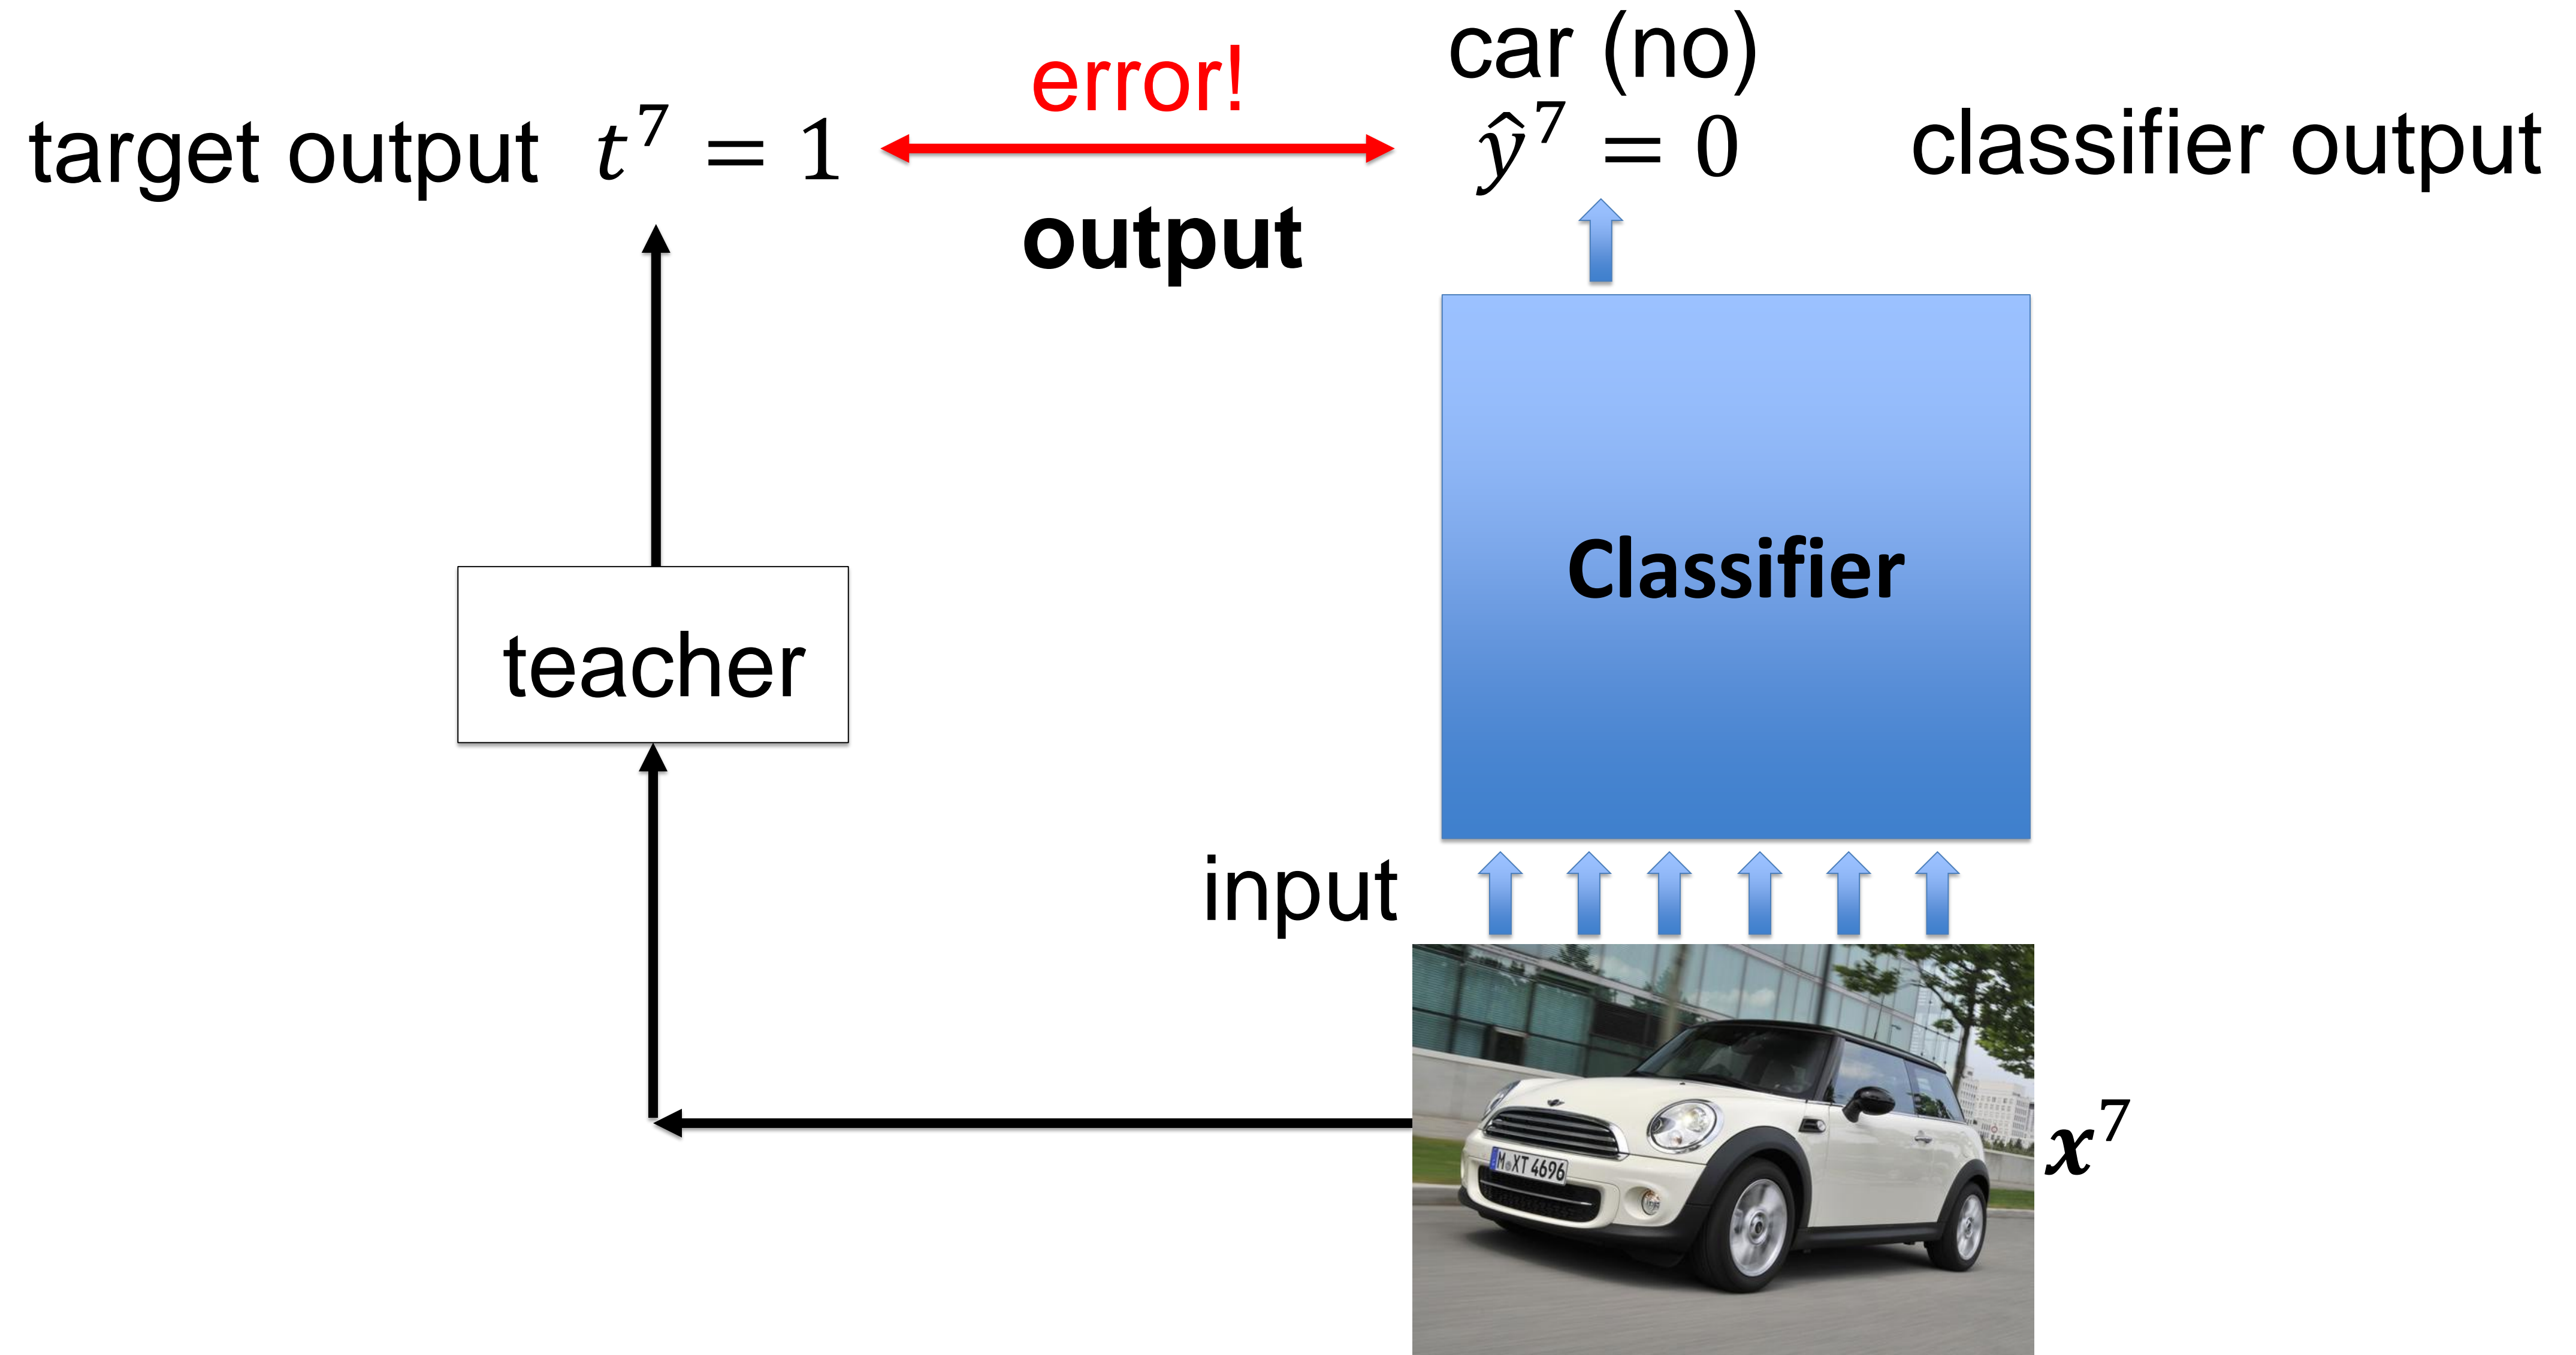
$P$  data points  $\{ (\mathbf{x}^\mu, t^\mu) , \quad 1 \leq \mu \leq P \};$

  
input    target output

$t^\mu = 1$     car =yes


$t^\mu = 0$     car =no

## 2. Data base for Supervised learning



## 2. Data base for Supervised learning

$P$  data points  $\{ (\mathbf{x}^\mu, t^\mu) , \quad 1 \leq \mu \leq P \};$



input    target output


for each data point  $\mathbf{x}^\mu$ , the classifier gives an output  $y^\mu$

**→ use errors  $\hat{y}^\mu \neq t^\mu$  for optimization of classifier**

**Remark:** for multi-class problems  $\mathbf{y}$  and  $\mathbf{t}$  are vectors

### 3. Single-Layer networks: simple perceptron

*Blackboard 3:*  
hyperplane






## *Blackboard 3:* hyperplane

### 3. Single-Layer networks: simple perceptron

$$\hat{y}^{\mu} = 0.5[1 + \text{sgn}(\sum_k w_k x_k - \vartheta)]$$

output

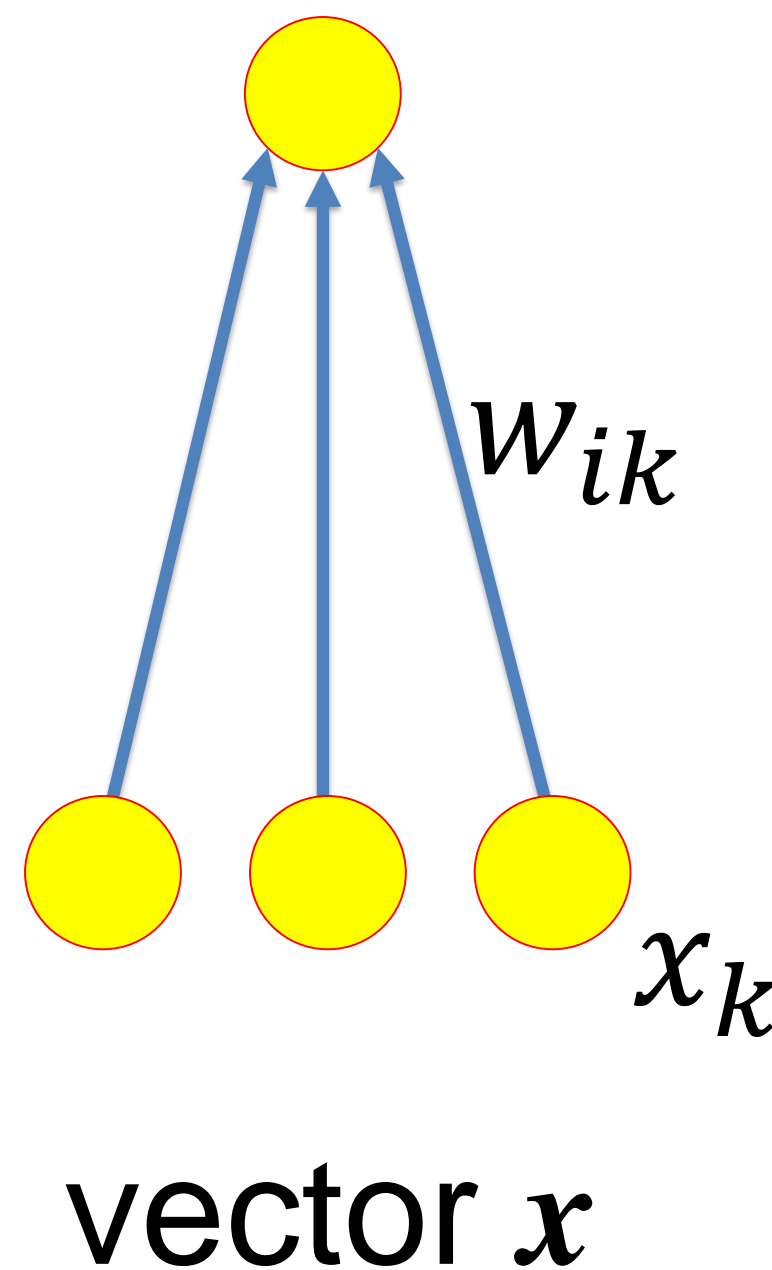


input


vector  $x$

### 3. Single-Layer networks: simple perceptron

$$\hat{y} = 0.5[1 + \text{sgn}(\sum_k w_k x_k - \vartheta)]$$




$$d(x) = \sum_k w_k x_k - \vartheta = 0$$




### 3. remove threshold: add a constant input

$$d(\mathbf{x}) = \sum_{k=1}^N w_k x_k - \vartheta = 0$$



$$d(\mathbf{x}) = \sum_{k=1}^{N+1} w_k x_k = 0$$





### 3. Single-Layer networks: simple perceptron

#### **a simple perceptron**

- can only solve linearly separable problems
- imposes a separating hyperplane
- for  $\vartheta = 0$  hyperplane goes through origin
- threshold parameter  $\vartheta$  can be removed by adding an input dimension  
→ in  **$N+1$**  dimensions hyperplane always goes through origin
- We can **adapt the weight vector** to the problem

#### 4. Perceptron algorithm: turn weight vector (in $N+1$ dim.)

$$\text{hyperplane: } d(\mathbf{x}) = \sum_{k=1}^{N+1} w_k x_k = \mathbf{w}^T \mathbf{x} = 0$$





## 4. Perceptron algorithm: turn weight vector

*Blackboard 4:*

geometry of perceptron algo

$$\Delta \mathbf{w} \sim \mathbf{x}^\mu$$

**Perceptron algo (in  $N+1$  dimensions):**

- set  $\mu = 1$
- (1) cycle many times through patterns

- choose pattern  $\mu$
- calculate output

$$\hat{y}^\mu = 0.5[1 + \text{sgn}(\mathbf{w}^T \mathbf{x}^\mu)]$$

- update by

$$\Delta \mathbf{w} = \gamma[t^\mu - \hat{y}^\mu]\mathbf{x}^\mu$$


- iterate  $\mu \leftarrow (\mu + 1) \bmod P$ , back to (1)
- (2) stop if no changes for all  $P$  patterns

*Blackboard 4:*  
geometry of the perc. algo

## 4. Perceptron algorithm: theorem

If the problem is linearly separable, the perceptron algorithm converges in a finite number of steps.

Proof: in many books, e.g.,  
Bishop, 1995,  
*Neural Networks for Pattern Recognition*



# Quiz: Perceptron algorithm



The **input vector has  $N$  dimensions** and we apply a perceptron algorithm.

- [ ] A change of parameters corresponds always to a rotation of the separating hyperplane in  $N$  dimensions.
- [ ] A change of parameters corresponds always to a rotation of the separating hyperplane in  $N+1$  dimensions.
- [ ] An increase of the length of the weight vector implies an increase of the distance of the hyperplane from the origin in  $N$  dimensions.
- [ ] An increase of the length of the weight vector implies that the hyperplane does not change in  $N$  dimensions
- [ ] An increase of the length of the weight vector implies that the hyperplane does not change in  $N+1$  dimensions

## 5. Sigmoidal output unit

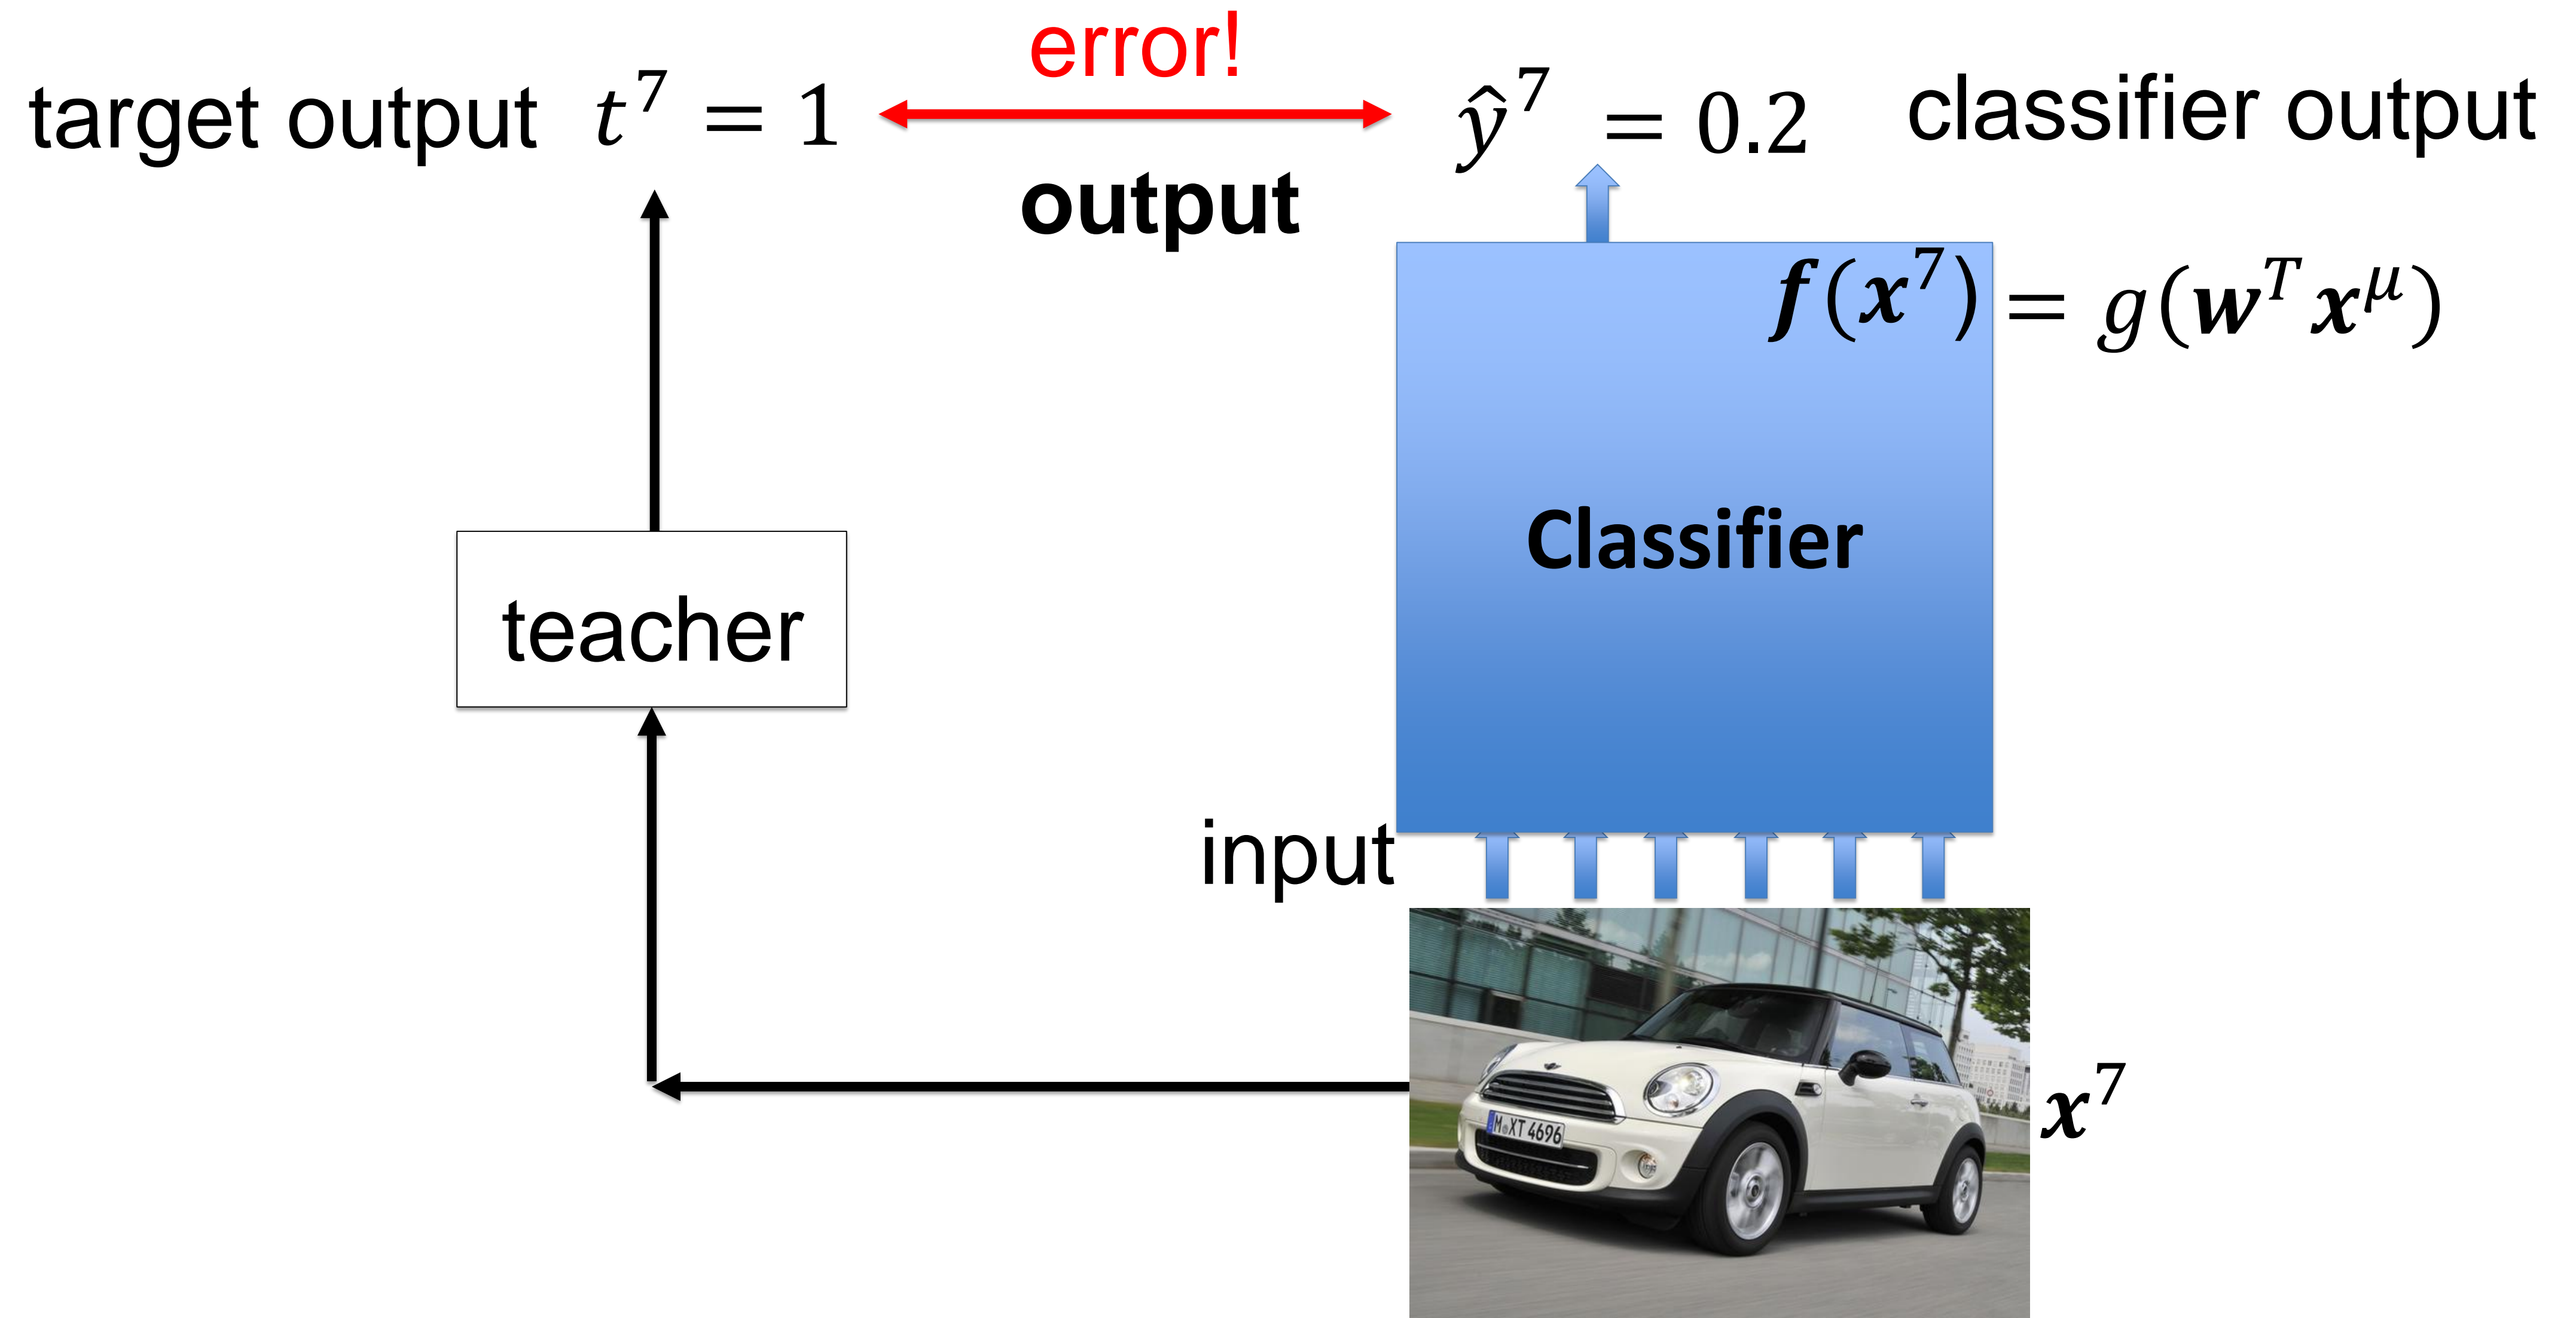
$$\hat{y}^{\mu} = g(\mathbf{w}^T \mathbf{x}^{\mu}) = g\left(\sum_{k=1}^{N+1} w_k x_k^{\mu}\right)$$

$$g(a) = \frac{\exp(a)}{1 + \exp(a)}$$






## 5. Supervised learning with sigmoidal output



## 5. Supervised learning with sigmoidal output


define error

$$E(\mathbf{w}) = \frac{1}{2} \sum_{\mu=1}^P \left[ t^{\mu} - \hat{y}^{\mu} \right]^2$$




gradient descent

$$w_k = -\gamma \frac{dE}{dw_k}$$



$$\hat{y}^{\mu} = g(\mathbf{w}^T \mathbf{x}^{\mu})$$




## 6. gradient descent

### Quadratic error

$$E(\mathbf{w}) = \frac{1}{2} \sum_{\mu=1}^P [t^{\mu} - \hat{y}^{\mu}]^2$$

gradient descent


$$w_k = -\gamma \frac{dE}{dw_k}$$



*Exercise 1.1 now:*

- calculate gradient (1 pattern)
- geometrical interpretation?

$$\hat{y}^{\mu} = g(\mathbf{w}^T \mathbf{x}^{\mu})$$




*Exercise 1.1 now:*

- *calculate gradient (1 pattern)*
- *geometrical interpretation?*

## 6. Gradient descent algorithm

$$\Delta \mathbf{w} = \gamma \delta(\mu) \mathbf{x}^\mu$$

- stepsize depends on (signed) output mismatch  $\delta(\mu)$  for this data point
- change implemented even if 'correctly classified
- compare with perceptron algorithm





## **Summary for today:**

- understand classification as a geometrical problem
- discriminant function of classification
- linear versus nonlinear discriminant function
- perceptron algorithm
- gradient descent for simple perceptrons

## Reading for this week:

**Bishop**, Ch. 4.1.7 of

*Pattern recognition and Machine Learning*

or

**Bishop**, Ch. 3.1-3.5 of

*Neural networks for pattern recognition*

## Motivational background reading:

Silver et al. 2017, Archive

*Mastering Chess and Shogi by Self-Play with a  
General Reinforcement Learning Algorithm*

**Goodfellow et al.**, Ch. 1 of

*Deep Learning*

